

Unsupervised Learning for Text

Joy Rimchala
IDEA DS&A
Intro to Data Science Series
2015.11.05

Motivation

- AI can help process human language to “understand” and “communicate” with human — elusive, long term
- Unsupervised NLP as a better tools — short term, more pragmatic
 - improve text search
 - information extraction, topic finding/discovery
 - sentiment analysis for marketing or campaign
 - automated/assisted machine translation
 - complex question answering

What's covered in this talk . . .

- Features from text (counts & frequencies) ~ 5 mins
 - word count (bag-of-word) and TF-IDF
 - word co-occurrence matrix
- Topic modeling ~ 15-20 mins
 - Latent Dirichlet Allocation (LDA)
- Vector Representation of Text ~ 10-15 mins
 - word2vec

Simple Feature Generation:

Bag of Word representation (BOW)

- Word counts - bag of word representation
 - Example codes in Notebook [1]-[3]
 - for each document,
 - split words by delimiter (usually white space)[1]
 - (additional pre-processing: removing stop words, stemming, lemmatizing)[1]
 - count word frequencies[3]

BOW in Python version I

```
file_name = './ted_mini/art_positive/5.ted'
delim = " "
with open(file_name, 'r') as f:
    dat = f.readlines()
    dat = map(lambda x: x.replace("\n", "").lower().split(delim), dat)
```

"""what i want to talk
this is actually quite
and so cars , as art ,
now at this point you "

['what', 'want', 'to', 'talk'
s', 'actually', 'quite', 'mea
'on', 'the', 'totem', 'pole',
f', 'it', 'and', 'cars', 'are
t', 'into', 'the', 'aesthetic

| | word | count |
|---|--------|-------|
| 0 | young | 639 |
| 1 | york | 638 |
| 2 | yelled | 637 |
| 3 | year | 636 |
| 4 | wrist | 635 |

```
d = pd.DataFrame(list(chain(*dat))).rename(columns={0: 'word'})
d['count'] = 1
df = d.groupby('word').agg('sum').sort('count', ascending=False)
```

BOW in Python version II

```
from sklearn.feature_extraction.text import CountVectorizer
cv = CountVectorizer()
vec = cv.fit(dat)
```

"""what i want to talk
this is actually quite
and so cars , as art ,
now at this point you "

['what', 'want', 'to', 'talk'
s', 'actually', 'quite', 'mea
'on', 'the', 'totem', 'pole',
f', 'it', 'and', 'cars', 'are
t', 'into', 'the', 'aesthetic

| | word | count |
|---|--------|-------|
| 0 | young | 639 |
| 1 | york | 638 |
| 2 | yelled | 637 |
| 3 | year | 636 |
| 4 | wrist | 635 |

```
df = pd.DataFrame(
    vec.vocabulary_.items())
    .rename(columns={0: 'word', 1: 'count'})
    .sort('count', ascending=0)
)
```

Simple Feature Generation:

Term Freq-Inverse Document Freq (TF-IDF)

- Term Frequency (TF): Normalized word counts:
 - Per document count of #times word appears in a document normalized by total number of words in document (i.e. per document BOW)
- Document Frequency (DF)
 - #documents containing a token normalized by total number of documents
- Inverse Document Frequency (IDF)
 - $1/DF$

Computing TF-IDF

- Corpus with 3 documents:

d₁: I like data science and data discovery. (7)

d₂: I think data science requires data exploration and machine learning (10)

d₃: I apply machine learning to data for science discovery (9)

- Corpus:

**{I, like, data, science, and, discovery, think, require,
exploration, machine, learning, apply, to, for}**

- Term Frequency (TF):

d₁: {I: 1/7, like: 1/7, data: 2/7...}

d₂: {I: 1/10, think: 1/10, data: 2/10, ...}

d₃: {I: 1/9, apply: 1/9, machine: 1/9, ...}

- Document Frequency (DF):

{I: 3/3, like: 1/3, data: 3/3, science: 3/3, and: 2/3, discovery: 2/3 , ...}

TF-IDF in Python

Poorman's version

```
def tf(t, doc):  
    if type(doc) == list:  
        return float(doc.count(t))/float(len(doc))  
    elif type(doc) == str:  
        doc = text_to_bagofwords(doc)  
        return float(doc.count(t))/float(len(doc))  
    else:  
        return NaN  
  
def idf(t, corpus):  
    df = 0  
    for doc in corpus:  
        if doc.count(t) > 0:  
            df += 1  
    return float(len(corpus))/float(df + 0.01)  
  
def tfidf(t, doc, corpus):  
    return tf(t, doc)*idf(t, corpus)
```

Off the shelf package

```
from sklearn.feature_extraction.text import TfidfVectorizer  
vec = TfidfVectorizer()  
tfidf = vec.fit_transform(dat)
```

What's topic modeling ?

- Topic modeling
 - discovers main themes that pervade collection of documents.
 - organizes the collection according to the discovered themes.
 - is unsupervised.
- Topic modeling algorithms can be adapted to many kinds of data including:
 - genetic data
 - images
 - social networks

Why topic modeling ?

- Get major themes or topics in text
 - Huge amount of documents
 - Want to know what's going on but can't read them all
- Unsupervised
- Simple way to analyze unlabeled text

Quick Review I

Multinomial distribution

- Given an observed sequence of die rolling what's the probability of an observed set of die outcome
- Die rolling is “generated” by a “hidden (Multinomial) process”
- The probability of this outcome set is described by Multinomial distribution

For topic modeling

- Three step dice rolling
 - per document topic distribution
 - topic die given topic distribution
 - word choice die given topic

Quick Review II

Chain Rule and Bayes' Rule

Chain Rule: $Pr(A \text{ and } B) = Pr(A|B) \cdot Pr(A)$

Bayes' Rule:

$$\underset{\text{posterior}}{Pr(B|A)} = \frac{\underset{\text{likelihood}}{Pr(A|B)} \cdot \underset{\text{prior}}{Pr(A)}}{\underset{\text{evidence}}{Pr(B)}}$$

For die rolling

- A = observed data (set of observed outcome)
- B = model parameters (probability of getting each face)

For topic modeling

- A = observed data (set of observed topic outcome)
- B = model parameters (probability of getting each topic per document, of getting word per topic, of getting topic proportion per document)

Quick Review III

Posterior Inference & Conjugate Prior

- Posterior Inference:
 - inferring model parameters from observed likelihood and prior
 - Likelihood: $Pr(A|B)$
 - Prob. of observing set of dice rolling outcome A given that it is generated by multinomial process with parameter set B
 - Prior: $Pr(A)$
 - Underlying knowledge about model parameters
 - if not know used “unbiased prior”
- Conjugate Prior:
 - posterior distribution is from the same family as prior time posterior
 - make the math for posterior inference easier
 - Conjugate prior of Multinomial distribution is **Dirichlet distribution**

Topic modeling algorithm:

Latent Dirichlet Allocation (LDA)

- Invented by David Blei (Toronto), Andrew Ng (Stanford), and Michael I. Jordan (Berkeley) as improvement to LSI and LSA
- Assume documents exhibit multiple topics
 - document \sim multinomial distribution over topics
 - Dirichlet distribution is a conjugate prior of multinomial
- Hidden topics in documents are generated by a 3-level Multinomial process
 - process is described by topic weights
 - topic is a distribution over words
 - words have different level of “membership” to a topic
 - documents consists of multiple topics in different proportion

LDA as three step nested Multinomial Process

$$\underbrace{Pr(B|A)}_{\text{posterior}} = \frac{\overbrace{Pr(A|B)}^{\text{likelihood}} \cdot \overbrace{Pr(A)}^{\text{prior}}}{\underbrace{Pr(B)}_{\text{evidence}}}$$

Word choosing die

- A = observed word frequencies in document
- B = model parameters (probability of getting each number for fair die)

Topic choosing die

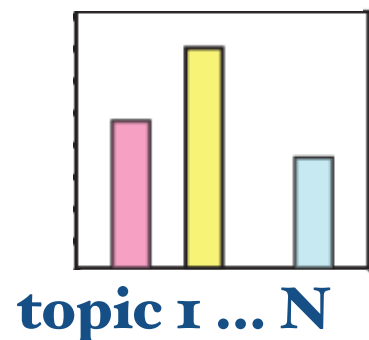
- A = observed topic distributions in document
- B = model parameters (probability of getting each topic given per topic distribution)

LDA assumes documents are generated by (hidden) Dirichlet Process

1. choose one of the distributions over words to generate each document

3. Repeat for all documents in collections

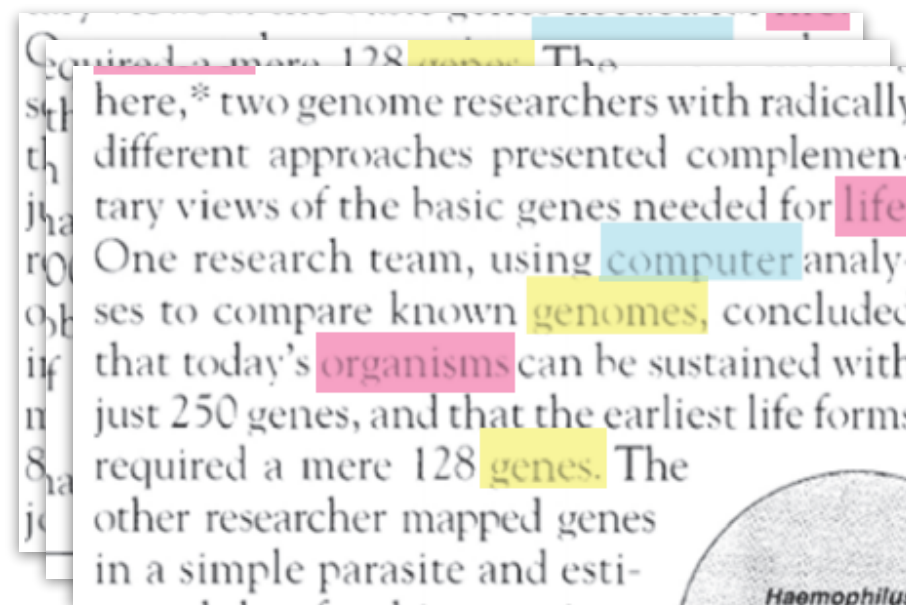
collection with topic distribution



collection of topics each with distribution over words

chosen topics

observed documents



discovered topic distribution over words

| | |
|---------|------|
| gene | 0.04 |
| dna | 0.02 |
| genetic | 0.01 |

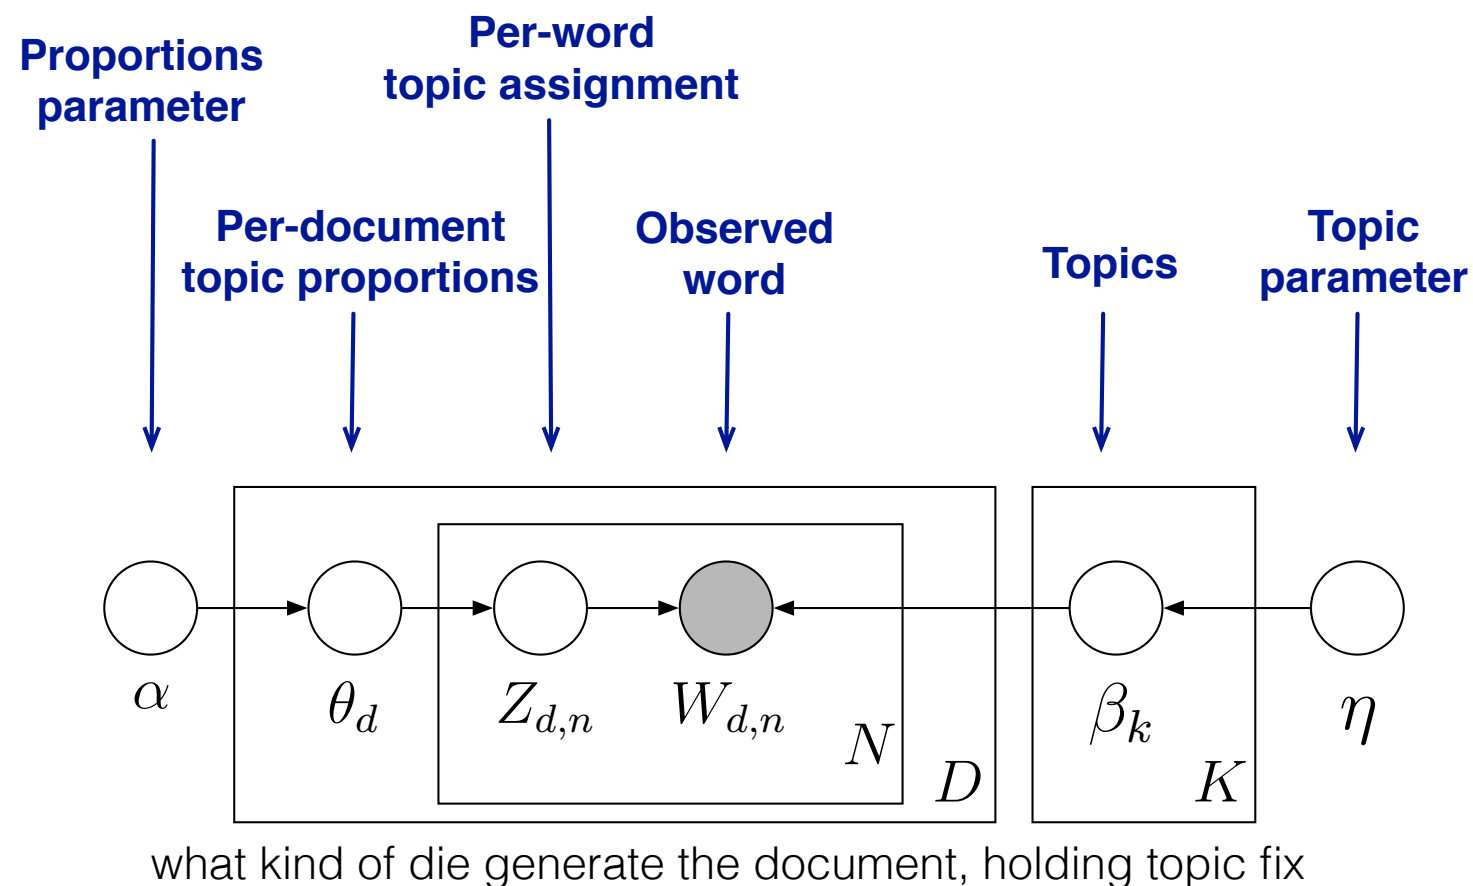
| | |
|----------|------|
| life | 0.02 |
| evolve | 0.01 |
| organism | 0.01 |

| | |
|----------|------|
| data | 0.02 |
| number | 0.02 |
| computer | 0.01 |

- In reality, observed documents consisting of words
- LDA takes all documents and infers underlying hidden distribution (structure) of topics

2. Choose word from one of the topics (blue, yellow, pink) and look up probability of chosen word in that topic

LDA as graphical model



LDA: reverse the generative process

Draw each topic for all topics 1, ..., K

For each document:

Draw topic proportion θ_d

For each word

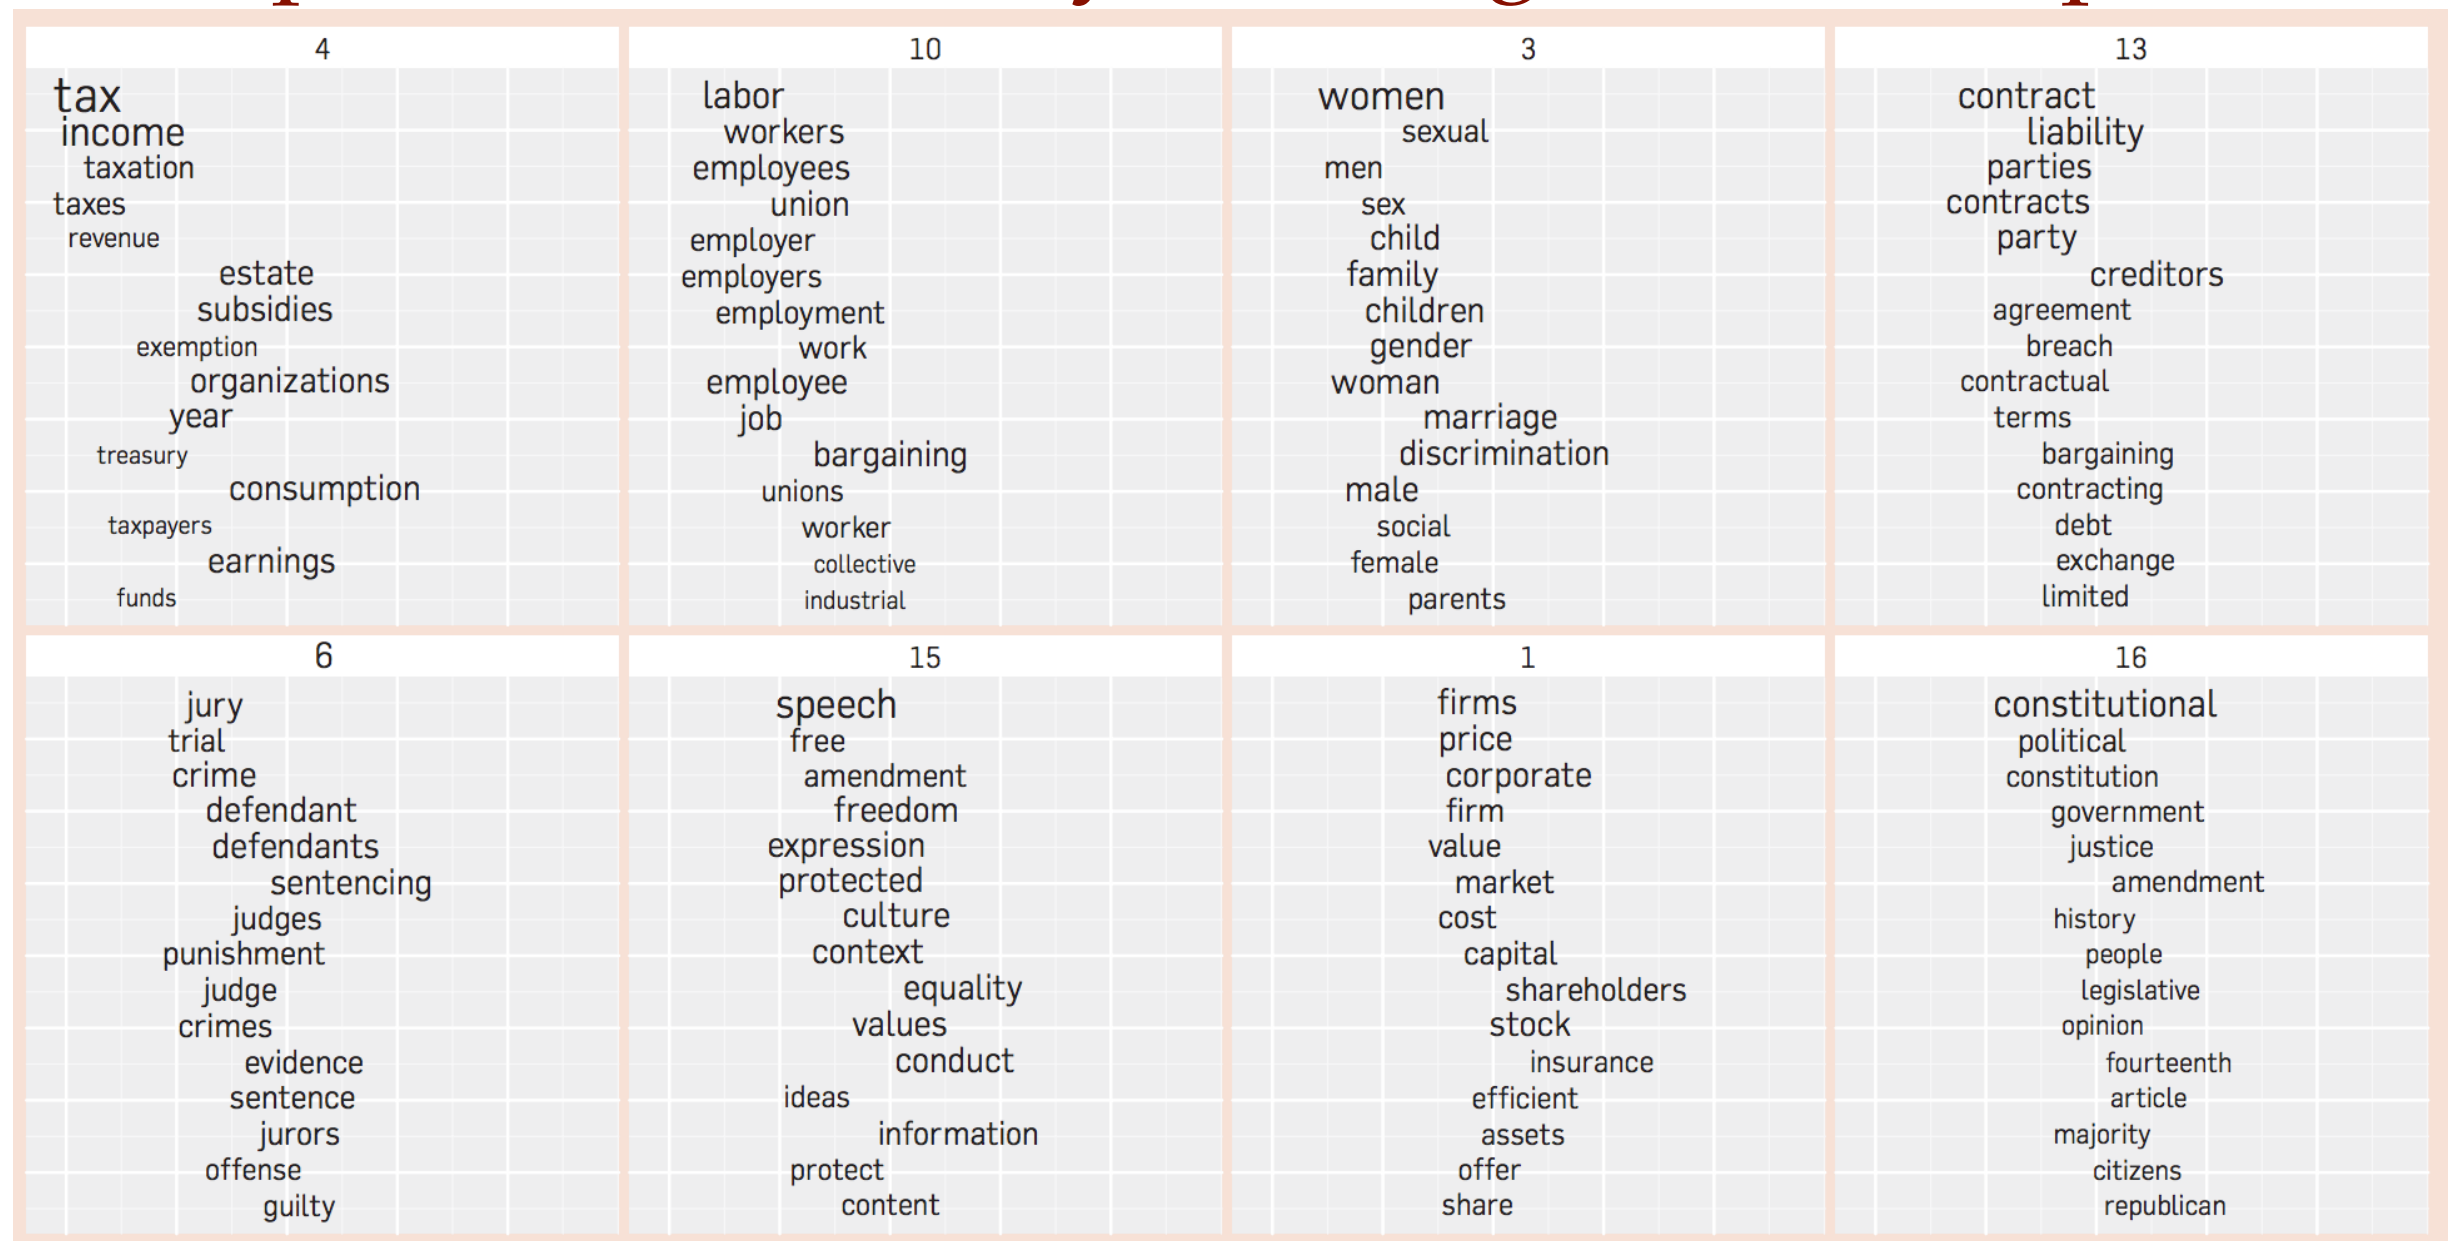
Draw topic assignment $Z_{d,n}$

Draw word $W_{d,n}$

β_k

Example LDA Results:

Topic auto-discovery from legal cases corpus



topic is a distribution over words

words have different level of “membership” to a topic

documents consists of multiple topics in different proportion

Results Taken from:

<http://www.cs.princeton.edu/~blei/papers/Blei2012.pdf>

LDA Input Computation

Co-occurrence Matrix by Hand

d_1 : I like data science and data discovery. (7)

d_2 : I think data science requires data exploration and machine learning (10)

d_3 : I apply machine learning to data for science discovery (9)

| | I | like | data | science | and | discovery | think | require | exploration | machine | learning | apply | to | for |
|-------------|---|------|------|---------|-----|-----------|-------|---------|-------------|---------|----------|-------|----|-----|
| I | 0 | 1 | 3 | 3 | 2 | 2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| like | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| data | 3 | 1 | 0 | 3 | 2 | 2 | 1 | 1 | 1 | 2 | 2 | 1 | 1 | 1 |
| science | 3 | 1 | 3 | 0 | 2 | 1 | 1 | 1 | 1 | 2 | 2 | 1 | 1 | 1 |
| and | 2 | 1 | 2 | 2 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| discovery | 2 | 1 | 2 | 1 | 1 | 0 | 1 | 0 | 0 | 1 | 1 | 1 | 1 | 1 |
| think | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 0 |
| require | 1 | 0 | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 1 | 0 | 0 | 0 |
| exploration | 1 | 0 | 1 | 1 | 1 | 0 | 1 | 1 | 0 | 1 | 1 | 0 | 0 | 0 |
| machine | 1 | 0 | 2 | 2 | 1 | 1 | 1 | 1 | 1 | 0 | 3 | 1 | 1 | 1 |
| learning | 1 | 0 | 2 | 2 | 1 | 1 | 1 | 1 | 1 | 3 | 0 | 1 | 1 | 1 |
| apply | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | 0 | 1 | 1 |
| to | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 |
| for | 1 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 0 |

LDA in Python lda

Reuters News corpus

```
import lda
model = lda.LDA(n_topics=20, n_iter=1500, random_state=1)
model.fit(X)
ntop=10
for i, topic_dist in enumerate(model.topic_word_[:ntop]):
    topic_words = np.array(vocab)[np.argsort(topic_dist)][:-(ntop+1):-1]
    print('Topic {}: {}'.format(i, ' '.join(topic_words)))
```

Topic 0: british churchill sale million major letters west
Topic 1: church government political country state people
Topic 2: elvis king fans presley life concert young death
Topic 3: yeltsin russian russia president kremlin moscow m
Topic 4: pope vatican paul john surgery hospital pontiff r
Topic 5: family funeral police miami versace cunanan city
Topic 6: simpson former years court president wife south c
Topic 7: order mother successor election nuns church nirma
Topic 8: charles prince diana royal king queen parker bowl
Topic 9: film french france against bardot paris poster an

LDA in Python gensim

TED corpus

```
from gensim.models.ldamodel import LdaModel
lda = LdaModel(corpus=artsci_corpus, id2word=dictionary,
               num_topics=100, update_every=1, chunksize=100, passes=1)
```

```
## Example topic
lda.show_topic(11)
```

```
[ (0.015640414022064685, 'earth'),
  (0.012160429080861469, 'stars'),
  (0.011924733515645118, 'like'),
  (0.011199539194868741, 'universe'),
  (0.010963916617258741, 'atoms'),
  (0.0095423565216509812, 'century'),
  (0.0085391585978015876, 'einstein'),
  (0.0079199531795733497, 'billion'),
  (0.0074580017159042661, 'planets'),
  (0.0073421488411868473, 'small')]
```

```
## Example topic
lda.show_topic(15)
```

```
[ (0.019069935795224396, 'ants'),
  (0.012175520975191553, 'nest'),
  (0.009606874068946494, 'ant'),
  (0.0095467406982767469, 'people'),
  (0.0092058530459285474, 'colony'),
  (0.0072401546937930416, 'like'),
  (0.0069742669948207143, 'workers'),
  (0.0062076249103784975, 'years'),
  (0.0060125350934175431, 'city'),
  (0.0054315333587530582, 'world')]
```

An Efficient discrete representation

Word2Vec

“You shall know a word by the company it keeps”

(J. R. Firth 1957: 11)

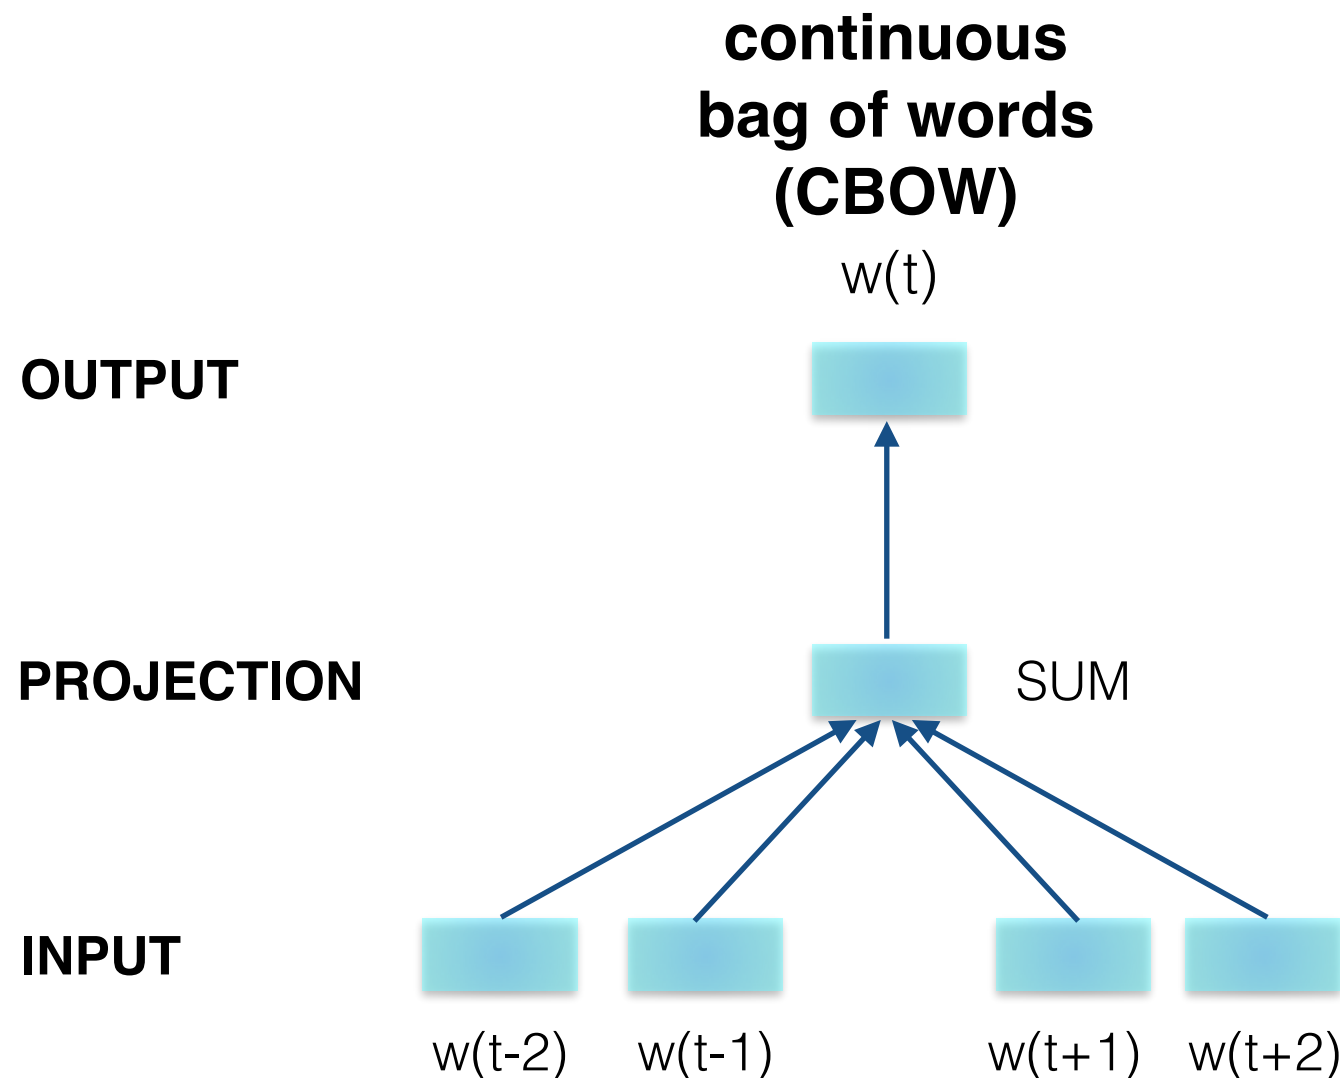
- Invented by Tomas Mikolov implemented in C with Python binding by Radim Řehůřek
- Assume that words that occur together share some semantic relationship
- How to make neighbors represent word ?
 - Window based word-word co-occurrence matrix
 - Predict surrounding words of every word in a window of length c

government debt problems turning into banking crises as has happened in
saying that Europe needs unified banking regulation to replace the hodgepodge

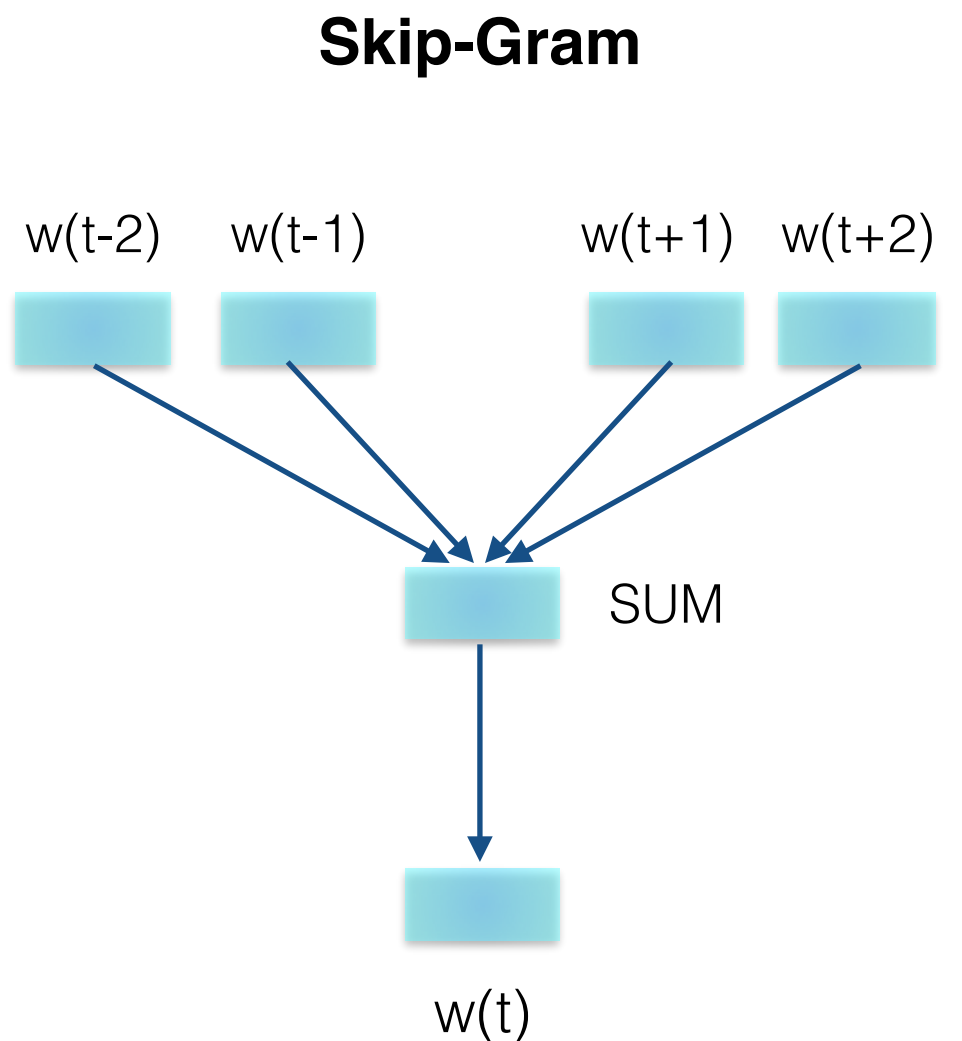
↖ These words will represent *banking* ↗

From Words to Vectors:

the two well known architectures



CBOW predicts the current word (inner vector) based on the surrounding words (outer vector, context)



Skip-Gram predicts surrounding words based on center word

From Words to Vectors:

Word2Vec Optimization problem

- Objective:
 - Maximize likelihood of any context word given current center word
 - Do this for every word in the corpus

Skip-Gram objective function:

maximize log-likelihood of center word
given surrounding words within a
specified window

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

Skip-Gram Likelihood function

softmax of dot products between
center word (inner vector) and
surrounding words (outer vectors)

$$p(w_O | w_I) = \frac{\exp \left(v'_{w_O}{}^\top v_{w_I} \right)}{\sum_{w=1}^W \exp \left(v'_w{}^\top v_{w_I} \right)}$$

Emerging Semantic Relationship in Word2Vec

Example Results

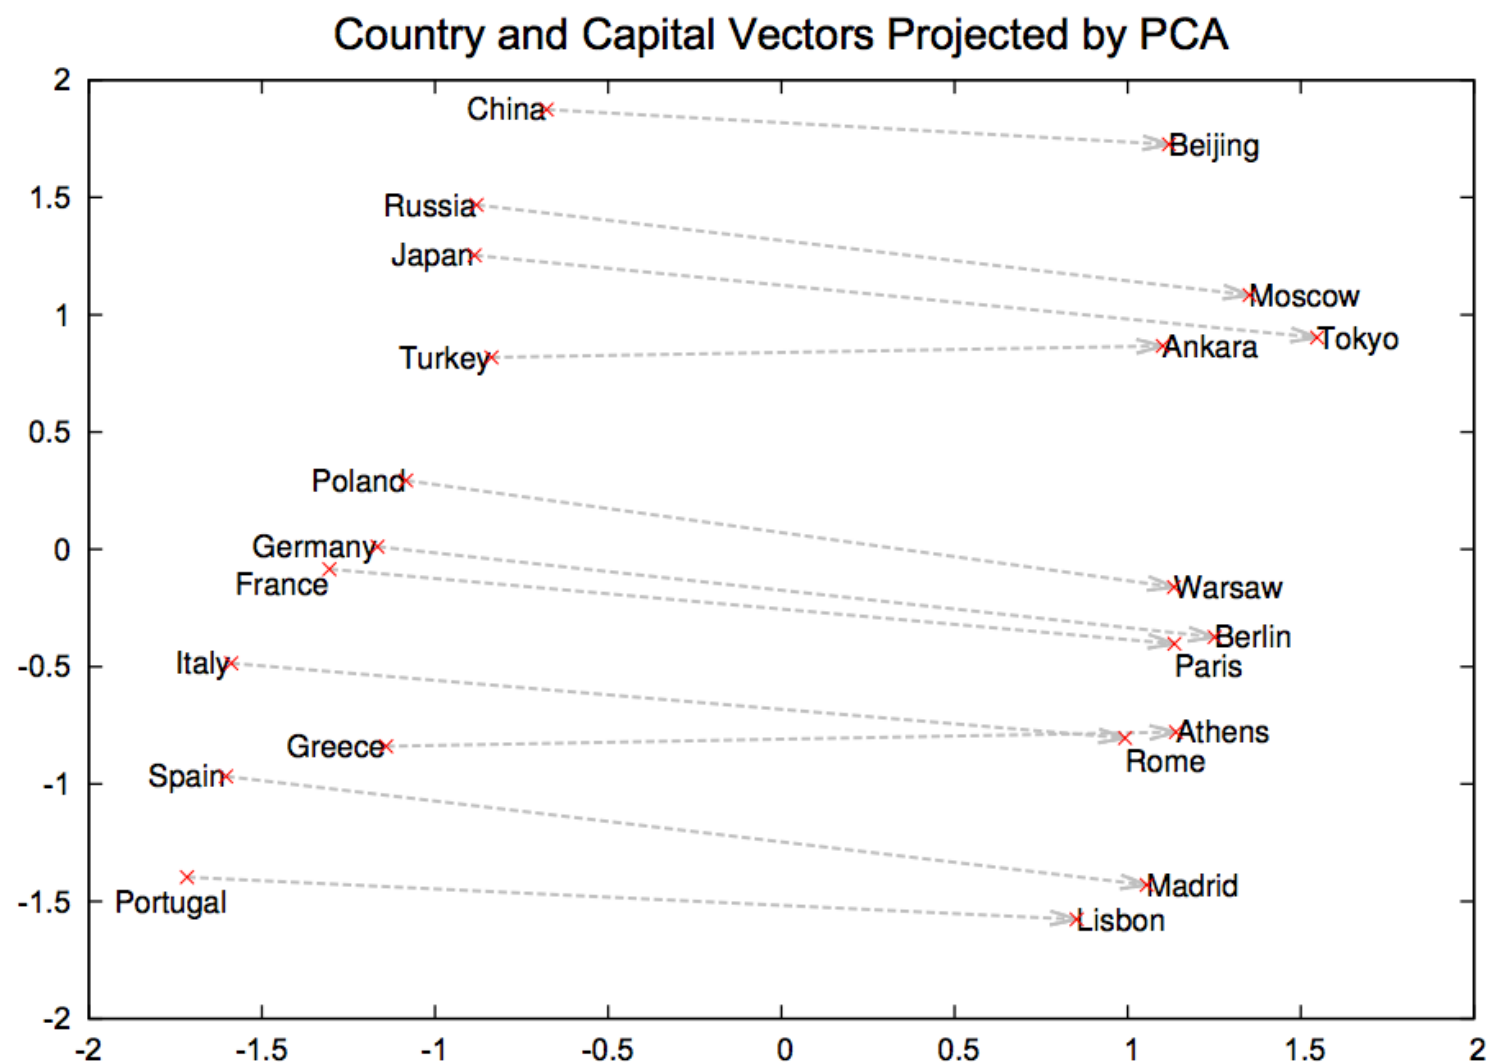
$X_{\text{king}} - X_{\text{man}} \sim X_{\text{queen}} - X_{\text{woman}}$

$X_{\text{apple}} - X_{\text{apples}} \sim X_{\text{car}} - X_{\text{cars}}$

$X_{\text{dog}} - X_{\text{dogs}} \sim X_{\text{family}} - X_{\text{families}}$

$X_{\text{shirt}} - X_{\text{clothing}} \sim X_{\text{chair}} - X_{\text{furniture}}$

Benchmark: Mikolov *et al* NIPS 2012



Skip-Gram Word2Vec in Python

```
from gensim.models.word2vec import Word2Vec
w2vmodel = Word2Vec(texts, size=100, window=5, min_count=5, workers=2)
w2vmodel.save(os.join.path(data_dir), 'artsci_positive_w2vmodel')
```

```
w2vmodel.similarity('man', 'woman')
```

```
0.71450582833706511
```

```
model.most_similar(positive=['cambridge', 'brain'], negative=['pittsburgh'], topn=1)
[('protein', 0.6212430000305176)]
```

```
model.doesnt_match("breakfast food neurons dinner".split())
```

```
'neurons'
```

Takeaways

- Simple features from text (counts & frequencies)
 - word count (bag-of-word) and TF-IDF: quick and easy to compute
 - word co-occurrence matrix: usually yield really sparse matrix
 - off-the-shelf in Python: **sklearn.feature_extraction.text** and **NLTK**
- Topic modeling: Latent Dirichlet Allocation (LDA)
 - Unsupervised method. Good for analyzing unlabeled text.
 - Computationally expensive to train, unclear measure of success (this is true for a lot of unsupervised learning algorithms)
 - off-the-shelf package in Python: **lda** and **gensim.model.Lda** (*not sklearn.Lda*)
- Vector Representation of Text (Skip-Gram word2vec)
 - Unsupervised method. Good for capturing semantic similarities across documents
 - off-the-shelf package in Python: **gensim.model.word2vec**